


Automated Benchmarking of Container Applications


Paulius Dilkas

1st August 2019

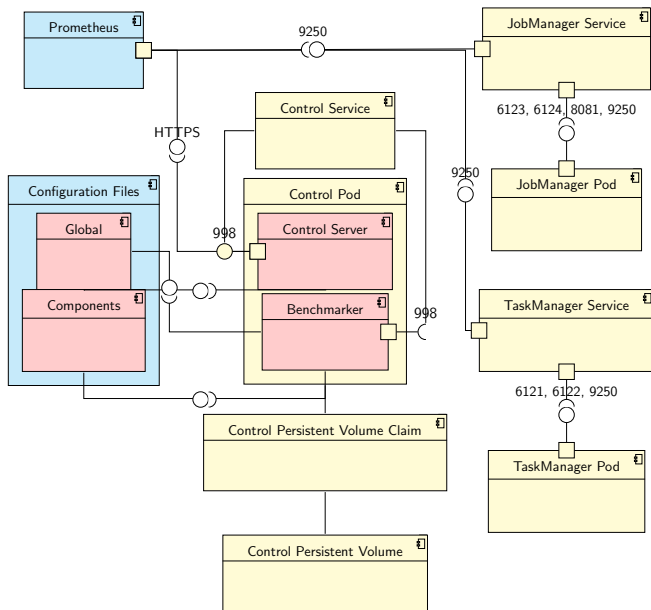
Main Ingredients



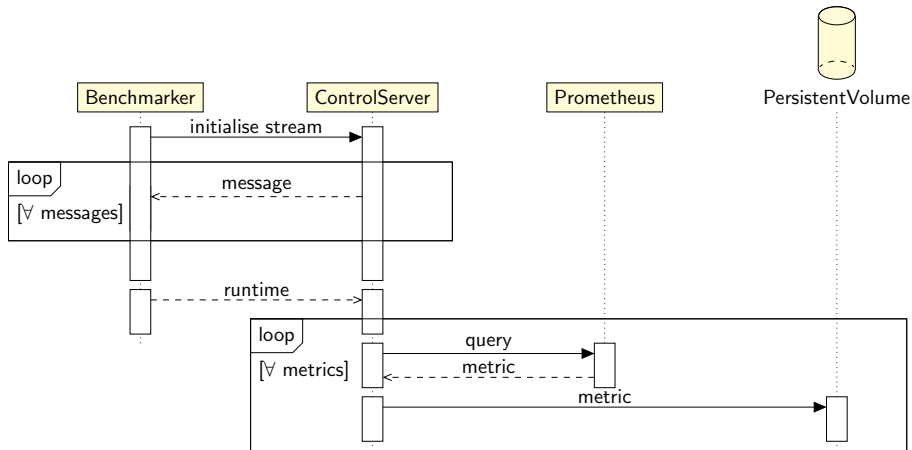
The General Idea




Deployment




Execution




Calibration




Evaluation: CPU




Evaluation: Memory (64 MiB)




Evaluation: Memory (128 MiB)




Evaluation: Memory (256 MiB)



Evaluation: Memory (512 MiB)



Using Maximum Value as a Predictor



Future Work

- Input/output simulation
- Complex usage patterns
- Automatically answering the question:
 - ▶ does this experiment show that the application could benefit from more resources?
- Complex component topologies
- More performance metrics (e.g. end-to-end latency)
 - ▶ and example applications